

Parallel Programming Tutorial - More on OpenMP

Bengisu Elis, M.Sc.

Philipp Czerner

Chair for Computer Architecture and Parallel Systems (Prof. Schulz)

Technical University Munich

29. Mai 2019



TUM Uhrenturm

Organizational

- **Q&A Sessions**

- Hands-on help with the assignments
- First session on 31st May
- Room : 01.06.020
- Fridays 08:15 - 09:45

- **Guest Lecture on next Monday**

- By Michael Klemm - Performance Engineer at Intel and CEO of the OpenMP ARB
- Inner workings of the OpenMP ARB
- (Mainly) How to utilize SIMD

Recap from last tutorial on OpenMP

Quiz; how to create a team of four threads to print their ids

```
1  #include <iostream>
2  #include<omp.h>
3
4  int main(){
5
6      int num_threads=4;
7      omp_set_num_threads(num_threads);
8
9      for (int i = 0; i < num_threads; i++)
10     {
11         std::cout << "My id is: "
12                 << omp_get_thread_num() << std::endl;
13     }
14 }
```

Quiz; how to create a team of four threads to print their ids

```
1  #include <iostream>
2  #include<omp.h>
3
4  int main(){
5
6      int num_threads=4;
7      omp_set_num_threads(num_threads);
8
9      for (int i = 0; i < num_threads; i++)
10     {
11         std::cout << "My id is: "
12                 << omp_get_thread_num() << std::endl;
13     }
14 }
```

./example1

My id is: 0
My id is: 0
My id is: 0
My id is: 0

Quiz (Cont.)

```

1  #include <iostream>
2  #include<omp.h>
3
4  int main(){
5
6      int num_threads=4;
7      omp_set_num_threads(num_threads);
8
9      #pragma omp for
10     for (int i = 0; i < num_threads; i++)
11     {
12         std::cout << "My id is: "
13                 << omp_get_thread_num() << std::endl;
14     }
15 }
```

Quiz (Cont.)

```

1  #include <iostream>
2  #include<omp.h>
3
4  int main(){
5
6      int num_threads=4;
7      omp_set_num_threads(num_threads);
8
9      #pragma omp for
10     for (int i = 0; i < num_threads; i++)
11     {
12         std::cout << "My id is: "
13                 << omp_get_thread_num() << std::endl;
14     }
15 }
```

./example2

My id is: 0
 My id is: 0
 My id is: 0
 My id is: 0

Quiz (Cont.)

```

1  #include <iostream>
2  #include<omp.h>
3
4  int main(){
5
6      int num_threads=4;
7      omp_set_num_threads(num_threads);
8
9      #pragma omp parallel
10     {
11         for (int i = 0; i < num_threads; i++)
12         {
13             #pragma omp critical
14             std::cout << "My id is: "
15                     << omp_get_thread_num() << std::endl;
16         }
17     }
18 }
```


Quiz (Cont.)

```

1  #include <iostream>
2  #include<omp.h>
3
4  int main(){
5
6      int num_threads=4;
7      omp_set_num_threads(num_threads);
8
9      #pragma omp parallel
10     {
11         for (int i = 0; i < num_threads; i++)
12         {
13             #pragma omp critical
14             std::cout << "My id is: "
15                     << omp_get_thread_num() << std::endl;
16         }
17     }
18 }

```

./example3

My id is: 3
 My id is: 0
 My id is: 3
 My id is: 0
 My id is: 3
 My id is: 0
 My id is: 3
 My id is: 0
 My id is: 3
 My id is: 0
 My id is: 1
 My id is: 1
 My id is: 1
 My id is: 1
 My id is: 2
 My id is: 2
 My id is: 2
 My id is: 2

Quiz (Cont.)

```

1  #include <iostream>
2  #include<omp.h>
3
4  int main(){
5
6      int num_threads=4;
7      omp_set_num_threads(num_threads);
8
9      #pragma omp parallel
10     {
11         #pragma omp parallel for
12         for (int i = 0; i < num_threads; i++)
13         {
14             #pragma omp critical
15             std::cout << "My id is: "
16                     << omp_get_thread_num() << std::endl;
17         }
18     }
19 }
```

Quiz (Cont.)

```

1  #include <iostream>
2  #include<omp.h>
3
4  int main(){
5
6      int num_threads=4;
7      omp_set_num_threads(num_threads);
8
9      #pragma omp parallel
10     {
11         #pragma omp parallel for
12         for (int i = 0; i < num_threads; i++)
13         {
14             #pragma omp critical
15             std::cout << "My id is: "
16                     << omp_get_thread_num() << std::endl;
17         }
18     }
19 }
```

./example4

My id is: 0
 My id is: 0
 My id is: 0
 My id is: 0
 My id is: 0
 My id is: 0
 My id is: 0
 My id is: 0
 My id is: 0
 My id is: 0
 My id is: 0
 My id is: 0
 My id is: 0
 My id is: 0
 My id is: 0

Quiz (Cont.)

```

1  #include <iostream>
2  #include<omp.h>
3
4  int main(){
5
6      int num_threads=4;
7      omp_set_num_threads(num_threads);
8      omp_set_nested(1);
9
10     #pragma omp parallel
11     {
12         #pragma omp parallel for
13         for (int i = 0; i < num_threads; i++)
14         {
15             #pragma omp critical
16             std::cout << "My id is: "
17                     << omp_get_thread_num() << std::endl;
18         }
19     }
20 }
```

Quiz (Cont.)

```

1  #include <iostream>
2  #include<omp.h>
3
4  int main(){
5
6      int num_threads=4;
7      omp_set_num_threads(num_threads);
8      omp_set_nested(1);
9
10     #pragma omp parallel
11     {
12         #pragma omp parallel for
13         for (int i = 0; i < num_threads; i++)
14         {
15             #pragma omp critical
16             std::cout << "My id is: "
17                     << omp_get_thread_num() << std::endl;
18         }
19     }
20 }
```

./example5

My id is: 1
 My id is: 0
 My id is: 2
 My id is: 3
 My id is: 1
 My id is: 2
 My id is: 0
 My id is: 1
 My id is: 1
 My id is: 0
 My id is: 3
 My id is: 2
 My id is: 3
 My id is: 0
 My id is: 3
 My id is: 2

Quiz (Cont.)

```

1  #include <iostream>
2  #include<omp.h>
3
4  int main(){
5
6      int num_threads=4;
7      omp_set_num_threads(num_threads);
8
9      #pragma omp parallel
10     {
11         #pragma omp for
12         for (int i = 0; i < num_threads; i++)
13         {
14             #pragma omp critical
15             std::cout << "My id is: "
16                     << omp_get_thread_num() << std::endl;
17         }
18     }
19 }
```

./example6

My id is: 0
 My id is: 1
 My id is: 2
 My id is: 3

Quiz (Cont.)

```

1  #include <iostream>
2  #include<omp.h>
3
4  int main(){
5
6      int num_threads=4;
7      omp_set_num_threads(num_threads);
8
9      #pragma omp parallel for
10     for (int i = 0; i < num_threads; i++)
11     {
12         #pragma omp critical
13         std::cout << "My id is: "
14                 << omp_get_thread_num() << std::endl;
15     }
16 }
```

./example7

My id is: 2
 My id is: 0
 My id is: 1
 My id is: 3

OpenMP Sections

OpenMP Sections

```
#pragma omp sections <{clause, ...}>
{
    #pragma omp section
    <structured block>

    #pragma omp section
    <structured block>
}
```

- The sections directive contains a set of structured blocks that are executed by single threads of a team
- Each structured block is preceded by a section directive
- The scheduling of the sections is implementation defined
- There is an implicit barrier at the end of a sections directive (unless `nowait`)
- Clauses: `private`, `firstprivate`, `lastprivate`, `reduction(identifier)`, `nowait`

Nested Regions

```
// environmmnet variable to set nested parallelism
OMP_NESTED
// library function to set/get nested parallelism
int omp_set_nested( int nested )
int omp_get_nested( void )
// limits/returns the number of maximal nested active parallel regions
int omp_set_max_active_levels( int max_levels )
int omp_get_max_active_levels( void )
// returns the number of current nesting level
int omp_get_level( void )
```

- Parallel regions and parallel sections may be arbitrarily nested inside each other
- If nested parallelism is disabled (default), the newly created team of threads will consist only of the encountering thread

Hint

- Take care of oversubscription when using nested parallelism.

Example: Traverse a binary tree

```

1 struct node
2 {
3     struct node *left, *right;
4     int key;
5     node(int k):key(k){}
6 };
7
8 void traverse(struct node *p)
9 {
10     if (p->left != NULL)
11         traverse(p->left);
12
13     if (p->right != NULL)
14         traverse(p->right);
15
16     process(p);
17 }

```

```

1 void process(struct node *p){
2     usleep(1000000);
3     std::cout << "element with key: "
4               << p->key << " is processed"
5               << std::endl;
6 }
7
8 int main(int argc, char *argv[])
9 {
10     struct node *tree = new struct node(0);
11     tree->left = new struct node(1);
12     tree->right = new struct node(2);
13     tree->left->left = new struct node(3);
14     tree->left->right = new struct node(4);
15     tree->right->left = new struct node(5);
16     tree->right->right = new struct node(6);
17
18     traverse(tree);
19     return 0;
20 }

```

Example: Traverse a binary tree (Cont.)

```

1 void traverse(struct node *p)
2 {
3     #pragma omp parallel
4     {
5         #pragma omp sections
6         {
7
8             #pragma omp section
9             {
10                 if (p->left != NULL)
11                     traverse(p->left);
12             }
13
14             #pragma omp section
15             {
16                 if (p->right != NULL)
17                     traverse(p->right);
18             }
19         }
20     }
21     process(p);
22 }

```

```

1 void process(struct node *p){
2     usleep(1000000);
3     #pragma omp critical
4     std::cout << "element with key: "
5                 << p->key << " is processed"
6                 << std::endl;
7 }
8 int main(int argc, char *argv[])
9 {
10     struct node *tree = new struct node(0);
11     tree->left = new struct node(1);
12     tree->right = new struct node(2);
13     tree->left->left = new struct node(3);
14     tree->left->right = new struct node(4);
15     tree->right->left = new struct node(5);
16     tree->right->right = new struct node(6);
17
18     omp_set_nested(1);
19     omp_set_max_active_levels(2);
20
21     traverse(tree);
22     return 0;
23 }


```

OpenMP Tasks

OpenMP Tasks

Why Tasks?

- We don't always deal with simple for loops for parallelization
- We don't always deal with simple data structures like arrays
- Some times we don't know the length of the loops at compile time e.g., while loop
- Some times we deal with unknown number of parallel sections
- We need to deal with parallelization of recursive algorithms
- It is possible without tasks (OpenMP 3.0) but it is not pretty



Task semantics

Terminology

task A specific instance of executable code and its data environment and ICVs.

task region A region consisting of all code encountered during the execution of a task.

explicit task A task generated when a task construct is encountered.

implicit task A task generated by an implicit parallel region.

tied task A task that, when its task region is suspended, can be resumed only by the same thread.

untied task A task that, when its task region is suspended, can be resumed by any thread in the team.

undelayed task A task for which execution is not deferred with respect to its generating task region.

included task A task for which execution is sequentially included in the generating task region.

merged task A task for which the data environment is the same as that of its generating task region.

Task semantics (Cont.)

```
#pragma omp task <{clause, ...}>
<structured block>
```

- Defines an explicit task, generated from the associated structured block.
- The encountering thread may immediately execute the task or defer it.
- Deferred tasks may be executed by any thread of the team.
- Tasks may be nested, but the task region of the inner task is not part of the task region of the outer task.
- A thread that encounters a task scheduling point (TSP) within a task may temporarily suspend this task.
- By default a task is tied to a thread (unless clause `untied`).

Task syntax

```
#pragma omp task <{clause, ...}>  
<structured block>
```

Clauses (not exhaustive)

- `if (<scalar logical expression>)`
if false, an undeferred task is generated
- `final (<scalar logical expression>)`
if true, the generated task and all child tasks are included (sequentialized) tasks
- `default (private | firstprivate | shared | none)`
default is firstprivate for tasks
- `mergeable`
if the generated task is an undeferred or included task, the generation may generate a merged task
- `private, firstprivate, shared (<list>)`

Task Scheduling Points (TSPs)

`#pragma omp taskyield`

- Specifies that the current task can be suspended (implicit TSP)

`#pragma omp taskwait`

- Specifies a wait on the completion of child tasks of the current task (implicit TSP)

`#pragma omp taskgroup`

- Specifies a wait on the completion of child tasks of the current task and their descendant tasks (implicit TSP)

`int omp_set_dynamic(int dynamic_threads)`

- Enables or disables dynamic adjustment of number of threads available for tasks in subsequent parallel regions

Task Scheduling

Whenever a thread reaches a TSP, the implementation may perform a task switch, implied by the following locations:

- immediately following the generation of an explicit task
- after the completion of a task region
- in a taskyield region
- in a taskwait region
- at the end of a taskgroup region
- in an implicit or explicit barrier region
- ...

Example 1: Hello world using tasks

```
1  #include <iostream>
2  #include <omp.h>
3
4  int main(int argc, char *argv[])
5  {
6      #pragma omp parallel
7      {
8          #pragma omp task
9          std::cout << "Hello World from task"
10                 << std::endl;
11      }
12      return 0;
13 }
```

OMP_NUM_THREADS=4 ./example1

Hello World from task
Hello World from task
Hello World from task
Hello World from task

Example 2: Which threads execute the tasks

```
1  #include <iostream>
2  #include <omp.h>
3
4  int main(int argc, char *argv[])
5  {
6      #pragma omp parallel
7      {
8          #pragma omp task
9          {
10             #pragma omp critical
11             std::cout << "Hello World from task,\n"
12                          << "    executed by thread: "
13                          << omp_get_thread_num()
14                          << std::endl;
15         }
16     }
17     return 0;
18 }
```

OMP_NUM_THREADS=4 ./example2

Hello World from task, executed by thread: 0
Hello World from task, executed by thread: 3
Hello World from task, executed by thread: 2
Hello World from task, executed by thread: 1

or

Hello World from task, executed by thread: 0
Hello World from task, executed by thread: 1
Hello World from task, executed by thread: 2
Hello World from task, executed by thread: 0

or ...

Example 3: Using a single thread to create tasks

```

1  int main(int argc, char *argv[])
2  {
3      #pragma omp parallel
4      {
5          #pragma omp single
6          {
7              for (int t = 0; t < omp_get_num_threads(); t++)
8              {
9                  #pragma omp task
10                 {
11                     #pragma omp critical
12                     std::cout << "Hello World from task,\
13                               executed by thread: "
14                               << omp_get_thread_num()
15                               << std::endl;
16                 }
17             }
18         }
19     }
20     return 0;
21 }

```

OMP_NUM_THREADS=4 ./example3

Hello World from task, executed by thread: 2
Hello World from task, executed by thread: 1
Hello World from task, executed by thread: 2
Hello World from task, executed by thread: 0

- Only one thread creates the tasks
- Unlike the previous example where all threads created tasks
- Created tasks can be nested and are scheduled to be executed by the available threads

Example 4: List traversal

```
1 void process_element(int &elem){
2     usleep(1000000);
3     std::cout << elem << std::endl;
4 }
5
6 void traverse_list(std::forward_list<int> &l){
7     for (auto it = l.begin(); it != l.end() ; it++) {
8         process_element(*it);
9     }
10 }
11
12 int main(int argc, char *argv[])
13 {
14     std::forward_list<int> l;
15     l.assign({0,1,2,3,4,5,6,7,8,9});
16
17     traverse_list(l);
18
19     return 0;
20 }
```

time ./example4

0
1
2
3
4
5
6
7
8
9

real 0m10.006s

Example 4: List traversal (Cont.)

```

1 void process_element(int &elem){
2     usleep(1000000);
3     #pragma omp critical
4     std::cout << elem << std::endl;
5 }
6
7 void traverse_list(std::forward_list<int> &l){
8     #pragma omp parallel
9     {
10         #pragma omp single
11         for (auto it = l.begin(); it != l.end() ; it++) {
12             #pragma omp task
13             process_element(*it);
14         }
15     }
16 }

```

time OMP_NUM_THREADS=4
./example4

0
1
2
3
4
5
6
7
8
9

real 0m3.015s

Example 5: Fibonacci Number

```

1  int fib(int n) {
2      int i, j;
3
4      if (n < 2) return n;
5
6      i = fib(n - 1);
7      j = fib(n - 2);
8
9      return i + j;
10 }
```

```

1  int main(int argc, char** argv) {
2      int n = 30;
3
4      if(argc > 1)
5          n= atoi(argv[1]);
6
7      printf("fib(%d) = %d\n", n, fib(n));
8
9  }
```

Example 5: Fibonacci Number (Cont.)

```
1  int fib(int n) {
2      int i, j;
3
4      if (n < 2) return n;
5
6      #pragma omp task shared(i) firstprivate(n)
7      i = fib(n - 1);
8
9      #pragma omp task shared(j) firstprivate(n)
10     j = fib(n - 2);
11
12     #pragma omp taskwait
13     return i + j;
14 }
```

```
1  int main(int argc, char** argv) {
2      int n = 30;
3
4      if(argc > 1)
5          n= atoi(argv[1]);
6
7      omp_set_num_threads(4);
8
9      #pragma omp parallel shared(n)
10     {
11         #pragma omp single
12         printf("fib(%d) = %d\n", n, fib(n));
13     }
14 }
```

Example 5: Fibonacci Number, Runtime

```
$ time ./fib 35
fib(35) = 9227465

real    0m9.785s
user    0m25.933s
sys     0m0.000s
```

Example 5: Fibonacci Number, final task

```
1  #define T 30 // THRESHOLD
2
3  int fib(int n)
4  {
5      int i, j;
6
7      if (n < 2)
8          return n;
9
10     #pragma omp task shared(i) firstprivate(n) final(n < T)
11     i = fib(n - 1);
12
13     #pragma omp task shared(j) firstprivate(n) final(n < T)
14     j = fib(n - 2);
15
16     #pragma omp taskwait
17     return i + j;
18 }
```

Example 5: Fibonacci Number, Runtime Final (GCC)

```
$ time ./fib_final 35
fib(35) = 9227465

real    0m0.392s
user    0m0.800s
sys     0m0.000s
```

Other directives

```
#pragma omp single <{clause, ...}>
```

- The single directive specifies that the associated block is executed by only one thread (not necessarily the master)
- The other threads of the team wait at an implicit barrier at the end of the single construct (unless `nowait`)
- Clauses: `private`, `firstprivate`, `nowait`

```
#pragma omp master <{clause, ...}>
```

- Same as single, but the thread is solely executed by the master thread
- Clauses: `private`, `firstprivate`, `nowait`

Other directives (Cont.)

`#pragma omp critical [<name>]`

- Restricts the execution of the associated structured block to a single thread at a time
- An optional name may be used to identify the critical construct
- All critical constructs without a name use a default name

`#pragma omp barrier`

- Specifies an explicit barrier
- All threads of a team must execute the barrier region
- Includes an implicit task scheduling point

Other directives (Cont.)

```
#pragma omp atomic [read | write | update | capture]  
<expression>
```

Example

```
#pragma omp atomic write  
x = 41;  
#pragma omp atomic  
{  
    v = x;  
    x++;  
}
```

- Ensures that a specific storage location is accessed atomically
- The expression reads|writes|read-writes|(read-writes + updates other variable) the storage location
- The structured block has two consecutive expressions
- To avoid race conditions, all accesses to a shared storage location must be protected with an atomic construct

Assignment 3 Solution

Assignment 3 Solution

```
1 #pragma omp parallel for num_threads(num_thrds) private(h,w) collapse(3)
2 for (i = 0; i < newImageHeight; i++) {
3     for (j = 0; j < newImageWidth; j++) {
4         for (d = 0; d < 3; d++) {
5             for (h = i; h < i + filterHeight; h++) {
6                 for (w = j; w < j + filterWidth; w++) {
7                     newImage[d][i][j] += filter[h-i][w-j] * image[d][h][w];
8                 }
9             }
10        }
11    }
12 }
```

- Remember to make private data private
- collapse does not really give any speedup here
- Order of floating point operations matters!

Assignment 4

Assignment 4: familytree

You have 1 week time for this assignment

Family Tree Algorithm

- The given algorithm computes the IQ for all members in a family.
 - It recursively traverses all generations (child \rightarrow {mother, father}).
 - At the end, all geniuses ($\text{IQ} \geq 140$) are printed at the end.
-
- Parallelize the sequential family tree algorithm with OpenMP.
 - Try to optimize it / reduce the overhead for tasking.
 - The goal is a speedup of ≥ 10 .

Assignment 4: familytree_seq.c

```

1  #include "familytree.h"
2
3  int traverse(tree *node, int num_threads) {
4      if (node == NULL) return 0;
5
6      int father_iq, mother_iq;
7
8      father_iq = traverse(node->father, num_threads);
9      mother_iq = traverse(node->mother, num_threads);
10
11     node->IQ = compute_IQ(node->data, father_iq, mother_iq);
12     genius[node->id] = node->IQ;
13     return node->IQ;
14 }

```

Assignment 4: familytree with OpenMP - Provided Files

- Makefile
 - contains rules to build executables
 - available targets: parallel, sequential, unit_test, all (default), clean
 - 'mode=debug make [target]' to build debug version, use 'make clean' before
- main.c
 - main function - argument handling + call familytree algorithm
- familytree.h
 - Header file for familytree.h and familytree_*.c
- familytree.c
 - Defines the familytree logic
- ds.h / ds.c
 - Header and definition for the needed datastructures
- familytree_seq.c
 - Sequential version of traverse().
- student/familytree_par.c
 - Implement the parallel version in this file

Assignment 4: familytree with OpenMP - Provided Files (Cont.)

- vis.h / vis.c
 - The visualization component
- unit_test.c
 - The unit tests that execute both the serial and parallel version to compare results.